

## VIII. APPENDIX

Details of equal\_load.py file,

"""

copyright@ BITS Pilani

Bhavik Dhandhaly

Anand Wani

"""

"""

Threshold for this file has been calculated manually and tested using "tat.py" file

"""

```
from pox.core import core
```

```
import pox
```

```
log = core.getLogger("iplb")
```

```
from pox.lib.packet.ethernet import ethernet, ETHER_BROADCAST
```

```
from pox.lib.packet.ipv4 import ipv4
```

```
from pox.lib.packet.arp import arp
```

```
from pox.lib.addresses import IPAddr, EthAddr
```

```
from pox.lib.util import str_to_bool, dpid_to_str
```

```
import pox.openflow.libopenflow_01 as of
```

```
import time
```

```
import random
```

```
FLOW_IDLE_TIMEOUT = 10
```

```
FLOW_MEMORY_TIMEOUT = 60
```

```
ALPHA = 0
```

```
class MemoryEntry(object):
```

```
    """
```

Record for flows we are balancing

Table entries in the switch "remember" flows for a period of time, but rather than set their expirations to some long value (potentially leading to lots of rules for dead connections), we let them expire from the switch relatively quickly and remember them here in the controller for longer.

Another tactic would be to increase the timeouts on the switch and use

the Nicira extension which can match packets with FIN set to remove them when the connection closes.

"""

```
def __init__(self, server, first_packet, client_port):
    self.server = server
    self.first_packet = first_packet
    self.client_port = client_port
    self.refresh()
```

```
def refresh(self):
    self.timeout = time.time() + FLOW_MEMORY_TIMEOUT
```

```
@property
def is_expired(self):
    return time.time() > self.timeout
```

```
@property
def from_client_to_server(self):
    ethp = self.first_packet
    ipp = ethp.find('ipv4')
    tcpp = ethp.find('tcp')

    return ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
```

```
@property
def from_server_to_client(self):
    ethp = self.first_packet
    ipp = ethp.find('ipv4')
    tcpp = ethp.find('tcp')

    return self.server, ipp.srcip, tcpp.dstport, tcpp.srcport
```

```
class iplb(object):
    """
```

A simple IP load balancer

Give it a service\_ip and a list of server IP addresses. New TCP flows to service\_ip will be randomly redirected to one of the servers.

We probe the servers to see if they're alive by sending them ARPs.

"""

```
def __init__(self, connection, service_ip, servers=[]):
```

```

self.service_ip = IPAddr(service_ip)
self.servers = [IPAddr(a) for a in servers]
self.con = connection
self.mac = self.con.eth_addr
self.live_servers = {} # IP -> MAC,port

try:
    self.log = log.getChild(dpid_to_str(self.con.dpid))
except:
    # Be nice to Python 2.6 (ugh)
    self.log = log

self.outstanding_probes = {} # IP -> expire_time

# How quickly do we probe?
self.probe_cycle_time = 5

# How long do we wait for an ARP reply before we consider a server dead?
self.arp_timeout = 3

self.total_connection = {} # IP -> total connection
for ip in servers:
    self.total_connection[ip] = 0

# We remember where we directed flows so that if they start up again,
# we can send them to the same server if it's still up. Alternate
# approach: hashing.
self.memory = {} # (srcip,dstip,srcport,dstport) -> MemoryEntry

self._do_probe() # Kick off the probing

# As part of a gross hack, we now do this from elsewhere
# self.con.addListener(self)

def _do_expire(self):
    """
    Expire probes and "memorized" flows
    Each of these should only have a limited lifetime.
    """
    t = time.time()

    # Expire probes
    for ip, expire_at in self.outstanding_probes.items():

```

```

        if t > expire_at:
            self.outstanding_probes.pop(ip, None)
            if ip in self.live_servers:
                self.log.warn("Server %s down", ip)
                del self.live_servers[ip]

# Expire flow
memory = self.memory.copy()
self.memory.clear()
for key, val in memory.items():
    ip = key[0]
    if ip in self.live_servers and val.is_expired:
        # Decrease total connection for that server
        self.total_connection[ip] -= 1
    if not val.is_expired:
        self.memory[key] = val

# Show information
# self.log.debug("Jumlah koneksi pada server:")
# for item in self.total_connection:
#     self.log.debug("%s = %s", item, self.total_connection[item])

# Expire old flows
# c = len(self.memory)
# self.memory = {k: v for k, v in self.memory.items()
#                 if not v.is_expired}
# if len(self.memory) != c:
#     self.log.debug("Expired %i flows", c - len(self.memory))

def _do_probe(self):
    """
    Send an ARP to a server to see if it's still up
    """
    self._do_expire()

    server = self.servers.pop(0)
    self.servers.append(server)

    r = arp()
    r.hwtype = r.HW_TYPE_ETHERNET
    r.prototype = r.PROTO_TYPE_IP
    r.opcode = r.REQUEST

```

```

r.hwdst = ETHER_BROADCAST
r.protodst = server
r.hwsrc = self.mac
r.protosrc = self.service_ip
e = ethernet(type=ethernet.ARP_TYPE, src=self.mac,
             dst=ETHER_BROADCAST)
e.set_payload(r)
# self.log.debug("ARPing for %s", server)
msg = of.ofp_packet_out()
msg.data = e.pack()
msg.actions.append(of.ofp_action_output(port=of.OFPP_FLOOD))
msg.in_port = of.OFPP_NONE
self.con.send(msg)

self.outstanding_probes[server] = time.time() + self.arp_timeout

core.callDelayed(self._probe_wait_time, self._do_probe)

@property
def _probe_wait_time(self):
    """
    Time to wait between probes
    """
    r = self.probe_cycle_time / float(len(self.servers))
    r = max(.25, r) # Cap it at four per second
    return r

def _pick_server(self, key, inport):
    """
    Pick a server for a (hopefully) new connection
    """
    global ALPHA

    if len(self.total_connection) == 0:
        return self.live_servers.keys()[0]
    ipserver = self.total_connection.keys()[0]
    totalconns = self.total_connection[ipserver]

    ALPHA = (ALPHA + 1) % 33
    if ALPHA < 28:
        for x in self.total_connection:
            if str(x) == '10.0.0.1':
                ipserver = x

```

```

        self.log.debug("Best available Server: %s" % ipserver)
        return ipserver
    elif ALPHA < 31:
        for x in self.total_connection:
            if str(x) == '10.0.0.2':
                ipserver = x
                self.log.debug("Best available Server: %s" % ipserver)
                return ipserver
    else:
        for x in self.total_connection:
            if str(x) == '10.0.0.3':
                ipserver = x
                self.log.debug("Best available Server: %s" % ipserver)
                return ipserver

    """
    for x in self.total_connection:
        if self.total_connection[x] < totalconns:
            ipserver = x
            totalconns = self.total_connection[x]
    """

    #self.log.debug("Best available Server: %s" % ipserver)
    #return ipserver

    # if len(self.total_connection) == 0:
    #     return self.live_servers.keys()[0]
    # return min(self.total_connection, key=self.total_connection.get)

def _handle_PacketIn(self, event):
    inport = event.port
    packet = event.parsed

    def drop():
        if event.ofp.buffer_id is not None:
            # Kill the buffer
            msg = of.ofp_packet_out(data=event.ofp)
            self.con.send(msg)
        return None

    tcp = packet.find('tcp')
    if not tcp:
        arpp = packet.find('arp')

```

```

if arpp:
    # Handle replies to our server-liveness probes
    if arpp.opcode == arpp.REPLY:
        if arpp.protosrc in self.outstanding_probes:
            # A server is (still?) up; cool.
            del self.outstanding_probes[arpp.protosrc]
            if (self.live_servers.get(arpp.protosrc, (None, None))
                == (arpp.hwsrc, inport)):
                # Ah, nothing new here.
                pass
            else:
                # Ooh, new server.
                self.live_servers[arpp.protosrc] = arpp.hwsrc, inport
                self.log.info("Server %s up", arpp.protosrc)
        return

    # Not TCP and not ARP. Don't know what to do with this. Drop it.
    return drop()

# It's TCP.

ipp = packet.find('ipv4')

# Incoming packet from server
if ipp.srcip in self.servers:
    key = ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
    entry = self.memory.get(key)

    if entry is None:
        # We either didn't install it, or we forgot about it.
        self.log.debug("No client for %s", key)
        return drop()

    # Refresh time timeout and reinstall.
    entry.refresh()

    # self.log.debug("Install reverse flow for %s", key)

    # Install reverse table entry
    mac, port = self.live_servers[entry.server]

    actions = []
    actions.append(of.ofp_action_dl_addr.set_src(self.mac))

```

```

actions.append(of.ofp_action_nw_addr.set_src(self.service_ip))
actions.append(of.ofp_action_output(port=entry.client_port))
match = of.ofp_match.from_packet(packet, inport)

msg = of.ofp_flow_mod(command=of.OFPFC_ADD,
                      idle_timeout=FLOW_IDLE_TIMEOUT,
                      hard_timeout=of.OFP_FLOW_PERMANENT,
                      data=event.ofp,
                      actions=actions,
                      match=match)
self.con.send(msg)

# Incoming packet from client
elif ipp.dstip == self.service_ip:
    # Ah, it's for our service IP and needs to be load balanced

    # Do we already know this flow?
    key = ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
    entry = self.memory.get(key)
    if entry is None or entry.server not in self.live_servers:
        # Don't know it (hopefully it's new!)
        if len(self.live_servers) == 0:
            self.log.warn("No servers!")
            return drop()

        # Pick a server for this flow
        server = self._pick_server(key, inport)
        self.log.debug("Directing traffic to %s", server)
        entry = MemoryEntry(server, packet, inport)
        self.memory[entry.from_client_to_server] = entry
        self.memory[entry.from_server_to_client] = entry

        # Increase total connection for that server
        self.total_connection[server] += 1

    # Update timestamp
    entry.refresh()

    # Set up table entry towards selected server
    mac, port = self.live_servers[entry.server]

    actions = []
    actions.append(of.ofp_action_dl_addr.set_dst(mac))

```



```

actions.append(of.ofp_action_nw_addr.set_dst(entry.server))
actions.append(of.ofp_action_output(port=port))
match = of.ofp_match.from_packet(packet, inport)

msg = of.ofp_flow_mod(command=of.OFPFC_ADD,
                      idle_timeout=FLOW_IDLE_TIMEOUT,
                      hard_timeout=of.OFP_FLOW_PERMANENT,
                      data=event.ofp,
                      actions=actions,
                      match=match)
self.con.send(msg)

# Remember which DPID we're operating on (first one to connect)
_dpid = None

def launch(ip, servers):
    servers = servers.replace(", ", " ").split()
    servers = [IPAddr(x) for x in servers]
    ip = IPAddr(ip)

    # Boot up ARP Responder
    from proto.arp_responder import launch as arp_launch
    arp_launch(eat_packets=False, **{str(ip): True})
    import logging
    logging.getLogger("proto.arp_responder").setLevel(logging.WARN)

def _handle_ConnectionUp(event):
    global _dpid
    if _dpid is None:
        log.info("IP Load Balancer Ready.")
        core.registerNew(iplb, event.connection, IPAddr(ip), servers)
        _dpid = event.dpid

    if _dpid != event.dpid:
        log.warn("Ignoring switch %s", event.connection)
    else:
        log.info("Load Balancing on %s", event.connection)

    # Gross hack
    core.iplb.con = event.connection
    event.connection.addListeners(core.iplb)

```

```

core.openflow.addListenerByName("ConnectionUp", _handle_ConnectionUp)

weightedRR.py
"""
copyright@ BITS Pilani
Bhavik Dhandhaly
Anand Wani
"""

"""
Threshold for this file has been calculated manually and tested using "tat.py" file
"""

from pox.core import core
import pox

log = core.getLogger("iplb")

from pox.lib.packet.ethernet import ethernet, ETHER_BROADCAST
from pox.lib.packet.ipv4 import ipv4
from pox.lib.packet.arp import arp
from pox.lib.addresses import IPAddr, EthAddr
from pox.lib.util import str_to_bool, dpid_to_str

import pox.openflow.libopenflow_01 as of

import time
import random

FLOW_IDLE_TIMEOUT = 10
FLOW_MEMORY_TIMEOUT = 60

ALPHA = 0

class MemoryEntry(object):
    """
    Record for flows we are balancing
    Table entries in the switch "remember" flows for a period of time, but
    rather than set their expirations to some long value (potentially leading
    to lots of rules for dead connections), we let them expire from the
    switch relatively quickly and remember them here in the controller for

```

longer.

Another tactic would be to increase the timeouts on the switch and use the Nicira extension which can match packets with FIN set to remove them when the connection closes.

"""

```
def __init__(self, server, first_packet, client_port):
```

```
    self.server = server
```

```
    self.first_packet = first_packet
```

```
    self.client_port = client_port
```

```
    self.refresh()
```

```
def refresh(self):
```

```
    self.timeout = time.time() + FLOW_MEMORY_TIMEOUT
```

```
@property
```

```
def is_expired(self):
```

```
    return time.time() > self.timeout
```

```
@property
```

```
def from_client_to_server(self):
```

```
    ethp = self.first_packet
```

```
    ipp = ethp.find('ipv4')
```

```
    tcpp = ethp.find('tcp')
```

```
    return ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
```

```
@property
```

```
def from_server_to_client(self):
```

```
    ethp = self.first_packet
```

```
    ipp = ethp.find('ipv4')
```

```
    tcpp = ethp.find('tcp')
```

```
    return self.server, ipp.srcip, tcpp.dstport, tcpp.srcport
```

```
class iplb(object):
```

```
    """
```

```
    A simple IP load balancer
```

```
    Give it a service_ip and a list of server IP addresses. New TCP flows to service_ip will be randomly redirected to one of the servers.
```

```
    We probe the servers to see if they're alive by sending them ARPs.
```

```
    """
```

```

def __init__(self, connection, service_ip, servers=[]):
    self.service_ip = IPAddr(service_ip)
    self.servers = [IPAddr(a) for a in servers]
    self.con = connection
    self.mac = self.con.eth_addr
    self.live_servers = {} # IP -> MAC,port

    try:
        self.log = log.getChild(dpid_to_str(self.con.dpid))
    except:
        # Be nice to Python 2.6 (ugh)
        self.log = log

    self.outstanding_probes = {} # IP -> expire_time

    # How quickly do we probe?
    self.probe_cycle_time = 5

    # How long do we wait for an ARP reply before we consider a server dead?
    self.arp_timeout = 3

    self.total_connection = {} # IP -> total connection
    for ip in servers:
        self.total_connection[ip] = 0

    # We remember where we directed flows so that if they start up again,
    # we can send them to the same server if it's still up. Alternate
    # approach: hashing.
    self.memory = {} # (srcip,dstip,srcport,dstport) -> MemoryEntry

    self._do_probe() # Kick off the probing

    # As part of a gross hack, we now do this from elsewhere
    # self.con.addListener(self)

def _do_expire(self):
    """
    Expire probes and "memorized" flows
    Each of these should only have a limited lifetime.
    """
    t = time.time()

```

```

# Expire probes
for ip, expire_at in self.outstanding_probes.items():
    if t > expire_at:
        self.outstanding_probes.pop(ip, None)
        if ip in self.live_servers:
            self.log.warn("Server %s down", ip)
            del self.live_servers[ip]

# Expire flow
memory = self.memory.copy()
self.memory.clear()
for key, val in memory.items():
    ip = key[0]
    if ip in self.live_servers and val.is_expired:
        # Decrease total connection for that server
        self.total_connection[ip] -= 1
    if not val.is_expired:
        self.memory[key] = val

# Show information
# self.log.debug("Jumlah koneksi pada server:")
# for item in self.total_connection:
#     self.log.debug("%s = %s", item, self.total_connection[item])

# Expire old flows
# c = len(self.memory)
# self.memory = {k: v for k, v in self.memory.items()
#                 if not v.is_expired}
# if len(self.memory) != c:
#     self.log.debug("Expired %i flows", c - len(self.memory))

def _do_probe(self):
    """
    Send an ARP to a server to see if it's still up
    """
    self._do_expire()

    server = self.servers.pop(0)
    self.servers.append(server)

    r = arp()
    r.hwtype = r.HW_TYPE_ETHERNET

```

```

r.prototype = r.PROTO_TYPE_IP
r.opcode = r.REQUEST
r.hwdst = ETHER_BROADCAST
r.protodst = server
r.hwsrc = self.mac
r.protosrc = self.service_ip
e = ethernet(type=ethernet.ARP_TYPE, src=self.mac,
             dst=ETHER_BROADCAST)
e.set_payload(r)
# self.log.debug("ARPing for %s", server)
msg = of.ofp_packet_out()
msg.data = e.pack()
msg.actions.append(of.ofp_action_output(port=of.OFPP_FLOOD))
msg.in_port = of.OFPP_NONE
self.con.send(msg)

```

```

self.outstanding_probes[server] = time.time() + self.arp_timeout

```

```

core.callDelayed(self._probe_wait_time, self._do_probe)

```

@property

```

def _probe_wait_time(self):

```

```

    """

```

```

    Time to wait between probes

```

```

    """

```

```

    r = self.probe_cycle_time / float(len(self.servers))

```

```

    r = max(.25, r) # Cap it at four per second

```

```

    return r

```

```

def _pick_server(self, key, inport):

```

```

    """

```

```

    Pick a server for a (hopefully) new connection

```

```

    """

```

```

    global ALPHA

```

```

    if len(self.total_connection) == 0:

```

```

        return self.live_servers.keys()[0]

```

```

    ipserver = self.total_connection.keys()[0]

```

```

    totalconns = self.total_connection[ipserver]

```

```

    ALPHA = (ALPHA + 1) % 33

```

```

    if ALPHA < 28:

```

```

        for x in self.total_connection:

```

```

        if str(x) == '10.0.0.1':
            ipserver = x
            self.log.debug("Best available Server: %s" % ipserver)
            return ipserver
    elif ALPHA < 31:
        for x in self.total_connection:
            if str(x) == '10.0.0.2':
                ipserver = x
                self.log.debug("Best available Server: %s" % ipserver)
                return ipserver
    else:
        for x in self.total_connection:
            if str(x) == '10.0.0.3':
                ipserver = x
                self.log.debug("Best available Server: %s" % ipserver)
                return ipserver

    """
    for x in self.total_connection:
        if self.total_connection[x] < totalconns:
            ipserver = x
            totalconns = self.total_connection[x]
    """
    #self.log.debug("Best available Server: %s" % ipserver)
    #return ipserver

    # if len(self.total_connection) == 0:
    #     return self.live_servers.keys()[0]
    # return min(self.total_connection, key=self.total_connection.get)

def _handle_PacketIn(self, event):
    inport = event.port
    packet = event.parsed

    def drop():
        if event.ofp.buffer_id is not None:
            # Kill the buffer
            msg = of.ofp_packet_out(data=event.ofp)
            self.con.send(msg)
        return None

    tcp = packet.find('tcp')

```

```

if not tcpp:
    arpp = packet.find('arp')
    if arpp:
        # Handle replies to our server-liveness probes
        if arpp.opcode == arpp.REPLY:
            if arpp.protosrc in self.outstanding_probes:
                # A server is (still?) up; cool.
                del self.outstanding_probes[arpp.protosrc]
                if (self.live_servers.get(arpp.protosrc, (None, None))
                    == (arpp.hwsrc, inport)):
                    # Ah, nothing new here.
                    pass
            else:
                # Ooh, new server.
                self.live_servers[arpp.protosrc] = arpp.hwsrc, inport
                self.log.info("Server %s up", arpp.protosrc)
        return

# Not TCP and not ARP. Don't know what to do with this. Drop it.
return drop()

# It's TCP.

ipp = packet.find('ipv4')

# Incoming packet from server
if ipp.srcip in self.servers:
    key = ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
    entry = self.memory.get(key)

    if entry is None:
        # We either didn't install it, or we forgot about it.
        self.log.debug("No client for %s", key)
        return drop()

    # Refresh time timeout and reinstall.
    entry.refresh()

    # self.log.debug("Install reverse flow for %s", key)

    # Install reverse table entry
    mac, port = self.live_servers[entry.server]

```



```

actions = []
actions.append(of.ofp_action_dl_addr.set_src(self.mac))
actions.append(of.ofp_action_nw_addr.set_src(self.service_ip))
actions.append(of.ofp_action_output(port=entry.client_port))
match = of.ofp_match.from_packet(packet, inport)

msg = of.ofp_flow_mod(command=of.OFPFC_ADD,
                      idle_timeout=FLOW_IDLE_TIMEOUT,
                      hard_timeout=of.OFP_FLOW_PERMANENT,
                      data=event.ofp,
                      actions=actions,
                      match=match)
self.con.send(msg)

# Incoming packet from client
elif ipp.dstip == self.service_ip:
    # Ah, it's for our service IP and needs to be load balanced

    # Do we already know this flow?
    key = ipp.srcip, ipp.dstip, tcpp.srcport, tcpp.dstport
    entry = self.memory.get(key)
    if entry is None or entry.server not in self.live_servers:
        # Don't know it (hopefully it's new!)
        if len(self.live_servers) == 0:
            self.log.warn("No servers!")
            return drop()

        # Pick a server for this flow
        server = self._pick_server(key, inport)
        self.log.debug("Directing traffic to %s", server)
        entry = MemoryEntry(server, packet, inport)
        self.memory[entry.from_client_to_server] = entry
        self.memory[entry.from_server_to_client] = entry

        # Increase total connection for that server
        self.total_connection[server] += 1

    # Update timestamp
    entry.refresh()

    # Set up table entry towards selected server
    mac, port = self.live_servers[entry.server]

```

```

        actions = []
        actions.append(of.ofp_action_dl_addr.set_dst(mac))
        actions.append(of.ofp_action_nw_addr.set_dst(entry.server))
        actions.append(of.ofp_action_output(port=port))
        match = of.ofp_match.from_packet(packet, inport)

        msg = of.ofp_flow_mod(command=of.OFPFC_ADD,
                               idle_timeout=FLOW_IDLE_TIMEOUT,
                               hard_timeout=of.OFP_FLOW_PERMANENT,
                               data=event.ofp,
                               actions=actions,
                               match=match)
        self.con.send(msg)

# Remember which DPID we're operating on (first one to connect)
_dpid = None

def launch(ip, servers):
    servers = servers.replace(", ", " ").split()
    servers = [IPAddr(x) for x in servers]
    ip = IPAddr(ip)

    # Boot up ARP Responder
    from proto.arp_responder import launch as arp_launch
    arp_launch(eat_packets=False, **{str(ip): True})
    import logging
    logging.getLogger("proto.arp_responder").setLevel(logging.WARN)

def _handle_ConnectionUp(event):
    global _dpid
    if _dpid is None:
        log.info("IP Load Balancer Ready.")
        core.registerNew(ip, event.connection, IPAddr(ip), servers)
        _dpid = event.dpid

    if _dpid != event.dpid:
        log.warn("Ignoring switch %s", event.connection)
    else:
        log.info("Load Balancing on %s", event.connection)

    # Gross hack

```

```
core.iplb.con = event.connection  
event.connection.addListener(core.iplb)
```

```
core.openflow.addListenerByName("ConnectionUp", _handle_ConnectionUp)
```